

AgentOS

An operating system for AI agents.

A complete reference: every component, why it exists, and where its limits are.

This manual assumes you know roughly what a process, a scheduler, and a system call are. It assumes nothing about AI agents, and nothing at all about this project. Everything else is explained as it comes up.

Read Part I to understand what AgentOS is and the one design decision everything else follows from. Read Part II for the kernel, component by component. Part III covers how agents actually run as operating-system processes; Part IV, how the system is used as a hosted service, including agents that are invented at runtime by a language model. Part V is evidence: what has been measured, and how. Part VI answers the questions a reader is most likely to be asked.

Contents

Part I — The idea ... 4

1. What AgentOS is
2. Why an operating system
3. Vocabulary
4. The one architectural decision

Part II — The kernel ... 7

5. Processes and the lifecycle
6. The scheduler
7. Syscalls
8. The event bus
9. The dependency graph
10. Human approval
11. Permissions and capabilities
12. Tool drivers
13. Memory
14. Model routing
15. The journal and crash recovery
16. The store

Part III — Running agents ... 16

17. Executors and isolation
18. The syscall transport

Part IV — The runtime as a service ... 18

19. The daemon and the HTTP API
20. Authentication
21. Agents invented at runtime
22. The safety model

Part V — Evidence ... 22

23. What has been measured
24. Compared with other frameworks

Part VI — Reference ... 25

25. Command line
26. File map
27. Questions and answers

Part I — The idea

1. What AgentOS is

AgentOS is a **runtime for AI agents**: a long-running program that owns agents the way an operating system owns processes. You hand it work; it decides when each agent runs, what each one is allowed to touch, what each one costs, and what happens when something crashes.

The distinction that matters is between a **library** and a **runtime**. Most agent tools today are libraries: you import them into your program, and your program is in charge. If your program stops, everything stops with it. AgentOS is a separate process that keeps running whether or not your program does — closer in shape to a database server than to a package you import.

Concretely, the system does four things a library cannot easily do. It **survives crashes** without repeating work, because every action an agent takes is written down before it proceeds. It **enforces limits** from outside the agent, so an agent cannot reach a tool it was not given even if it tries. It **accounts for everything** in one place, so the question “what did all of this cost?” has an answer across every application using it. And it lets agents **be created at runtime** by a language model, without giving up any of the above.

The shortest possible summary. Agents are processes. Everything an agent does goes through the kernel as a system call. Because every call is written down, a crash costs nothing. Because every call is checked, an agent can only do what it was permitted. Because every call is priced, you always know the bill.

2. Why an operating system

The analogy is not decoration; it is the design. A traditional operating system exists because programs cannot be trusted to share a machine politely. Left alone, one program would hog the processor, read another's memory, and corrupt the disk. The OS solves this by putting itself in the middle: programs stop touching the hardware directly and start asking the kernel for things.

Agent systems have the same problem with different resources. Instead of CPU and memory, agents contend over **language models** (which cost money per call), **tools** (which have real side effects — files written, emails sent, rows deleted), **human attention** (someone has to approve the risky things), and **each other** (one agent's output is another's input). Left alone, an agent framework has no answer for “who is allowed to spend money?” or “what happens if this dies halfway?”

So AgentOS applies the same move. Agents stop calling tools and models directly, and start asking a kernel. The table below is the mapping, and it is worth reading closely — most of the rest of this manual is one row of it at a time.

In a traditional OS	In AgentOS	What it manages
Process	Agent	A unit of work with an identity, a lifecycle, and a parent
System call	Syscall over a message boundary	The only way an agent affects anything
Scheduler	Scheduler with execution slots	Which agent runs next, and how many at once
Device driver	Tool driver	Files, shells, databases, the network
File permissions	Capability grants	What each agent is allowed to reach
Virtual memory	Memory manager	Six kinds of agent memory, private and shared

In a traditional OS	In AgentOS	What it manages
Journalling filesystem	Syscall journal	Surviving a crash without redoing work
CPU time accounting	Token and dollar ledger	What each agent spent
(no equivalent)	Human approval	A person as a first-class dependency

One row has no equivalent. Traditional operating systems have no notion of “wait for a human to approve this.” Agent systems need it constantly, because agents take actions people want a say in. AgentOS treats a waiting human as exactly the same kind of thing as a waiting tool call — see §10.

3. Vocabulary

These terms recur throughout. None is standard outside this project, so they are defined once here.

Term	Meaning
Agent	A unit of work the runtime manages. In code, a Python class with a <code>run()</code> method. At runtime, a process with a PID.
Kernel	The part of AgentOS that owns everything: the process table, the scheduler, permissions, the journal. Agents never hold a reference to it.
Syscall	A request from an agent to the kernel — “spawn this child”, “call this tool”, “wait for that event”. Travels as JSON.
Context (<code>ctx</code>)	The object an agent uses to make syscalls. It is the agent's <i>entire</i> view of the world.
Capability	The right to use one tool, named as a string: <code>filesystem</code> , <code>shell</code> , <code>sql</code> .
Grant	The act of giving an agent capabilities. A parent may only grant what it holds.
Driver	The code behind a capability. Owns retries, timeouts, rate limits, caching for that one tool.
Event	A named announcement (“ <code>MeasurementsReady</code> ”) that wakes whoever is waiting for it.
Journal	The record of every completed syscall and its result. The basis of crash recovery.
Checkpoint	A point a crashed agent can resume from. In AgentOS every completed syscall is one.
Daemon	The long-running server process that hosts the kernel and serves the HTTP API.
Task	A goal submitted to the daemon, plus the whole tree of agents created to satisfy it.

4. The one architectural decision

Everything in this system follows from a single rule, and it is worth understanding before anything else:

An agent never holds a reference to the kernel, to the process table, or to another agent. Its entire world is the Context object. Every call it makes crosses a boundary as plain JSON and comes back as plain JSON.

In ordinary Python this would be unusual. You would normally let an agent hold a reference to a tool object and call methods on it. AgentOS deliberately refuses. An agent that wants to read a file does not get a file object; it sends a message describing what it wants and waits for a message describing what happened.

```
# What an agent CAN do – every line is a message to the kernel
result = await ctx.request_tool("filesystem", "read", path="notes.txt")
pid    = await ctx.spawn(Researcher(topic="trees"), grant=["http"])
await ctx.publish("ResearchDone", findings=result)
answer = await ctx.request_model("fast", prompt="Summarise: ...")

# What an agent CANNOT do – there is no such attribute, by design
ctx.kernel          # does not exist
ctx.processes       # does not exist
other_agent.run()    # raises DirectInvocationError
```

This constraint is checked at runtime, not merely documented: a syscall carrying something unserializable (a function, an open socket, a live object) is rejected the moment it is made. That refusal seems pedantic until you notice what it buys.

What the rule pays for

- **Crash recovery.** Everything an agent does crosses the boundary as JSON, so the kernel can write down every reply. After a crash, the agent is re-run and the recorded replies are handed back instantly instead of re-executing. The tool does not fire twice; the model is not billed twice. (§15)
- **Real process isolation.** Anything that survives conversion to JSON survives a network socket. So agents can be moved into separate operating-system processes without a single line of agent code changing — the messages simply travel further. (§17)
- **Agents invented at runtime.** If an agent's identity is JSON parameters, then a language model can write those parameters, and the resulting agent is as real as a hand-written one. It journals, recovers and isolates identically. (§21)

Those three capabilities normally cost a rewrite each. Here they are consequences of one rule applied consistently from the beginning.

Part II — The kernel

The kernel is the part that owns everything. This part of the manual walks through it one component at a time. Each section says what the component is for, how it behaves, and where its edges are.

5. Processes and the lifecycle

Every agent the runtime knows about has an entry in the **process table**, exactly as a program does in an operating system. The entry holds a PID (a number, assigned in order), a name, a parent, a list of children, a status, a priority, what it is waiting on, which model it last used, which capabilities it holds, how much it has spent, and how many checkpoints it has passed.

That whole record is what `agent ps` prints:

\$ python -m agentos.cli ps										
PID	NAME	PARENT	CHILDREN	STATUS	PRIORITY	WAITING ON	MODEL	PERMS	CKPT	
1	Planner	-	2	Running	Normal	-	gpt-oss	filesystem	#11	
2	Surveyor	1	0	Waiting	Normal	tool filesystem	gpt-oss	filesystem	#8	
3	Analyst	1	0	Blocked	High	Senior Engineer	gpt-oss	-	#6	

The nine states

An agent is always in exactly one of nine states. The important thing is not the list but the fact that the transitions between them are **enforced**: an illegal move raises an error rather than quietly corrupting the table.

State	Meaning	Typically moves to
Ready	Runnable, waiting for the scheduler to grant it a slot	Running
Running	Currently executing agent code	Waiting, Sleeping, Blocked, Finished
Waiting	Blocked on a tool call, a model call, another agent, or an event	Ready
Sleeping	Blocked on a timer it asked for	Ready
Blocked	Waiting for a <i>human</i> to approve something	Ready
Checkpointing	Flushing a durable point to disk	Running
Suspended	Paused by an operator with <code>agent pause</code>	Ready
Finished	Completed and returned a result	— (terminal)
Failed	Raised an error, or was killed	Ready, but only via a retry

The most important edge in the whole system. When an agent is woken — its tool finished, its event fired, its human approved — it goes to **Ready**, not straight to **Running**. It rejoins the queue and waits for a slot like everyone else. That single detail is the difference between a scheduler and a callback: the runtime decides what runs next, not whoever happened to finish. A test fails if this ever regresses.

Failed is worth a note. It is terminal for the agent's run — the `alive` property counts it as dead — but the scheduler may restart a failed agent within a retry budget, so there is a real edge from Failed back to Ready. This is covered in §6.

6. The scheduler

The scheduler answers one question: which Ready agent runs next? It has three parts worth understanding — slots, policies, and the loop.

Execution slots

A slot is permission to run. If the runtime has four slots, at most four agents execute at a time regardless of how many exist; the rest sit in Ready. This is the same idea as a thread pool, and it exists for the same reason: unbounded concurrency is how systems fall over. Set it with `--slots`.

Policies

Policy	How it picks	Use when
fifo	Oldest Ready agent first	The default. Fair and predictable.
priority	High before Normal before Low, with an <i>ageing</i> guard so a starved agent eventually rises	Some work genuinely matters more
dependency	Whoever unblocks the most other agents	Deep dependency chains where one agent is a bottleneck

The dependency-aware policy is the one a normal OS cannot have. A CPU scheduler has no idea that six processes are all waiting on process 12. This kernel knows exactly that, because agents declare what they are waiting for (§9) — so it can run the agent that frees the most work.

The loop, and why it is event-driven

The kernel runs a loop: drain incoming syscalls, hand out slots, check for deadlock, repeat. The loop is **event-driven with a tick as its ceiling**, not a fixed poll. A syscall or a newly-runnable agent wakes it in microseconds; the tick only bounds how long it may doze when genuinely nothing is happening.

This detail was a real performance bug once, and the reason is instructive. The loop originally slept for the tick on every pass. But `asyncio.sleep()` cannot resolve below the operating system's timer granularity — about 15.6 milliseconds on Windows — so *every syscall* paid that penalty no matter how small the tick was set. Ticks below the quantum were an illusion. Making the loop wake on work rather than on a clock took a durable step from roughly 31ms to 3.5ms.

Retries

An agent that raises an exception can be restarted, up to a budget (`Kernel(retries=N)`, or per-agent). The restart is a real state transition — Failed back to Ready — so it appears in the log rather than happening behind the state machine's back.

The subtle part: a restarted agent *replays its journal minus the failed tail*. Everything that succeeded is handed back instantly; the call that failed is re-executed for real. Without that trimming, a retry would replay its own recorded failure forever and be a restart in name only. An agent that a human *killed* is never retried — a person said stop.

7. Syscalls

A syscall is the only way an agent affects anything. There are eleven, and they are the complete surface of the Context object. If a capability is not in this table, an agent does not have it.

Syscall	What it does	Blocks?
spawn	Create a child agent, optionally granting it capabilities and wiring its events	No
sleep	Give up the slot for N seconds; state becomes Sleeping	Yes
wait_all	Wait for a <i>set</i> of things: agents, events, a timer	Yes
log	Write a line to the kernel log	No
publish	Announce an event; the runtime wakes subscribers	No
subscribe	Register interest in an event type	No
memory	store / retrieve / share / delete, across six memory kinds	No
request_tool	Ask for a capability by name; the kernel checks and dispatches	Yes
request_model	Ask for a model by capability class, never by name	Yes
request_approval	Block until a human with a named role approves	Yes
checkpoint	Mark a durable point explicitly; passes through Checkpointing	No

“Blocks” means the agent gives up its execution slot and changes state. A non-blocking syscall is handled inside the kernel loop and the agent continues immediately. Either way the reply is journaled before the agent proceeds — which is what makes every completed syscall a checkpoint.

8. The event bus

Agents never call each other. There is no method by which one agent can invoke another — `Context` offers no way to even *name* another agent except by PID, and calling another agent’s `run()` directly raises `DirectInvocationError`. Instead they announce things, and the runtime decides who cares.

```
# The publisher does not know who is listening, or whether anyone is:
await ctx.publish("ResearchCompleted", topic="trees", findings=[...])

# Elsewhere, in an agent that has never heard of the publisher:
await ctx.subscribe("ResearchCompleted")
payload = await ctx.wait_event("ResearchCompleted")
```

The benefit is that adding a fourth subscriber to an event requires editing no existing agent. This is not a style guideline — it is enforced by the absence of any other mechanism.

Delivery is buffered, not broadcast

Each subscriber has its own queue per event type. An event that fires while a subscriber is busy still lands in that subscriber’s queue and waits. Without this buffer, “did I subscribe before you published?” becomes a race — and races in a scheduler are the bugs you never reproduce.

The eight kernel events

The kernel publishes these itself, so any agent can react to them without coordination: `AgentFinished`, `AgentFailed`, `ToolCompleted`, `HumanApproved`, `MemoryUpdated`, `ModelFinished`, `TimerExpired`, `FileCreated`. Applications may define any other name they like.

Is the event bus a database? No. It is plain in-memory data inside the kernel — a queue per subscriber and a history list. Delivery never touches disk. Events *are* also written to SQLite, along with who was subscribed and who consumed what, but that record exists purely so a restarted runtime can redeliver an event that was owed and never taken. That is recovery, not delivery.

9. The dependency graph

This is the piece that replaces workflow diagrams. Instead of writing “first A, then B, then C”, an agent declares what it is waiting *for*, and the scheduler wakes it when the last of those things resolves. Nobody wrote the order; the graph produced it.

```
# "Wake me when all three of these are true" – not a sequence, a set.
result = await ctx.wait_all(
    agents=[market_pid, legal_pid],      # two other agents must finish
    events=["HumanApproved"],            # and this must have fired
    timer=5,                             # and five seconds must have passed
)
```

A dependency can be another agent, an event, a timer, a running tool call, a running model call, or a human approval. They are all the same kind of node in the graph, which is why a human being slow is not a special case in the code.

Deadlock detection, in two flavours

- **Cycles are refused when requested.** If agent A waits on B and B already waits on A, the second wait does not hang — it raises immediately, naming the cycle, so the agent that closed the loop learns about it.
- **Stalls are detected.** If every surviving agent is Waiting, nobody is Sleeping, and no timer is pending, nothing can ever happen again. The runtime reports this rather than sitting there.
- **Unsatisfiable event waits fail fast.** If an agent waits on an event that *provably* nobody will publish, the wait is refused at once. And if the last agent that could have published it exits while someone is waiting, that waiter is failed then and there, with the cause named.

Why “provably” matters. A wait is only refused when every live agent has a declared event vocabulary (§21) and none of them lists it. If a single agent could still publish anything, the wait stands. Kernel events are always waitable. The system refuses only when it can prove the wait is hopeless — never on suspicion.

10. Human approval

Some actions want a person's sign-off. In most frameworks this is a callback or a confidence threshold. Here it is a kernel object — the same kind of thing as a tool call or a timer.

```
# The agent stops here. Its status becomes Blocked, waiting on the named role.
approval = await ctx.request_approval(
    role="Senior Engineer",
    reason="Production deployment",
)

# Meanwhile, from another terminal, possibly days later:
$ python -m agentos.cli approve 1 --as "Senior Engineer"
```

Approving with the wrong role is refused. An agent blocked on a human is not treated as a deadlock — the runtime keeps serving everyone else.

Why it survives restarts

The approval is written to durable storage, not held in memory. This has three consequences that a callback could not offer:

- Kill the runtime while an agent is blocked, start it again, and the re-run agent re-attaches to the *same* pending approval instead of asking a second time.
- A human can approve while the runtime is not even running. The next start-up sails straight through.
- Every grant also publishes a `HumanApproved` event, so other agents can react to it.

Why this is the hard part to fake. A human dependency that evaporates when the process restarts is not a kernel object — it is a callback with good marketing. Durability is what makes the difference observable.

11. Permissions and capabilities

An agent does not import tool libraries. It asks for a **capability** by name, and the kernel decides. The check happens *before* dispatch, in the kernel, so the application does not get a vote and the agent cannot route around it.

There are two ways an agent can hold a capability, and understanding the difference matters.

By name: the permission matrix

A JSON file maps agent names to capabilities. Deny is the default — an agent holds nothing it was not given. The file is re-read when it changes, so revoking a capability affects a *running* system: the next request after the edit is refused.

```
// .agentos/permissions.json
{
  "Finance": ["sql", "email"],
  "Coder":   ["filesystem", "python"],
  "**":      []                // everyone else gets nothing
}
```

By process: delegation and attenuation

The matrix works when you wrote the agents and know their names. It fails completely when agents are invented at runtime, because every one of them might be called the same thing. So capabilities can instead ride on the *process*, handed down by a parent:

```
# I hold {filesystem, http}. I may give a child any part of that — never more.
await ctx.spawn(Surveyor(...), grant=["filesystem"]) # fine
await ctx.spawn(Attacker(...), grant=["shell"])      # refused by the kernel
```

Attenuation is the whole security model. Because a parent can only pass on a subset of what it holds, the capability set given to the agent at the top of a task is the **ceiling for the entire tree** — however many layers of agents get invented underneath it, and whatever they are named. Start a task with `[filesystem]` and nothing in it can ever reach a shell. §22 reports what happened when a real language model was told to break exactly this.

Two details make it airtight rather than approximate. A per-process grant *overrides* the name matrix, so a deliberately narrowed child cannot be widened again by being named something privileged. And a grant dies with the process that held it.

12. Tool drivers

A driver is the code behind a capability — the equivalent of a device driver. It owns the messy parts once, so no agent has to: authentication, timeouts, rate limiting, retries, caching, and turning failures into something an agent can understand rather than a stack trace.

Capability	Operations	Notes
filesystem	read, write, list, exists	Sandboxed to a root directory. Publishes <code>FileCreated</code> .
shell	run	Runs a command. Real code execution — grant carefully.
python	run	Runs code in a fresh interpreter. Also real code execution.
sql	query, execute	SQLite. <code>query</code> is a read; <code>execute</code> writes.
http	get, post	Plain HTTP requests.
browser	open, get	Fetches a page and extracts readable text.

The shared discipline

Every driver inherits the same behaviour, so it is written once: a per-attempt timeout; a minimum interval between calls (rate limiting); retries for failures marked as transient; and an optional cache.

Caching is **opt-in per operation**, and the reason is worth stating: caching a write would be a correctness bug, not an optimisation. Each driver declares which of its operations are reads — `sql.query` yes, `sql.execute` never — and a cache is only used when a time-to-live is configured. The default is off.

A running tool call is a node in the dependency graph like anything else. The agent shows `Waiting on tool sql`, completion publishes `ToolCompleted`, and the woken agent rejoins the queue for a slot. Every dispatch is recorded, so the runtime can report which tools were used and how often they failed.

13. Memory

Six kinds of memory behind four verbs — `store`, `retrieve`, `share`, `delete`. The storage backend is invisible to agents; today it is SQLite, and swapping in Redis or a vector database would change one file and no agent.

Kind	Who can read it	How long it lives
working	Only the agent that wrote it	Freed when that agent exits
scratchpad	Same as working, by convention for rough notes	Freed when that agent exits
shared	Whoever the owner shared it with	The run
longterm	Any agent with the same <i>name</i>	Survives restarts
semantic	Same as longterm, plus similarity search	Survives restarts
episodic	The agent's own history; read-only	Written by the kernel

Shared memory is the only way agents pass state. Not by touching each other — there is no mechanism for that — but through the kernel, with an access list. This matters most for agents invented at runtime: an event payload is fine for a notification, but a dataset has to go through shared memory.

Two subtleties. **longterm** and **semantic** are keyed by agent *name*, not PID, which is why they survive restarts — a new agent with the same name inherits what the last one learned. And the semantic embedding is a deliberately humble stand-in (a hashed bag of words) that a real embedding model can replace without any agent changing, because agents only ever say `query=...`

14. Model routing

Agents ask for a **capability class** — “fast”, “reasoning” — and never for a model by name. Which actual model answers is a configuration decision, not an application one.

```
reply = await ctx.request_model("fast", prompt="Summarise: ...")
reply["text"], reply["model"], reply["cost"]
```

The routing table lists candidates in order. A candidate is skipped if it is unavailable — its API key is not set, or the prompt is longer than its context window — and one that fails at call time falls through to the next. The shipped chain is one open-weight model reachable four ways:

```
"fast": [
  {"provider": "openai", "base_url": "https://api.groq.com/openai/v1",
   "api_key_env": "GROQ_API_KEY", "model": "openai/gpt-oss-120b"},
  {"provider": "openai", "base_url": "https://openrouter.ai/api/v1",
   "api_key_env": "OPENROUTER_API_KEY", "model": "openai/gpt-oss-120b"},
  {"provider": "openai", "base_url": "http://localhost:11434/v1",
   "model": "gpt-oss:120b", "api_key_env": null},           // local Ollama
  {"provider": "mock", "model": "mock-fast"}               // always answers
]
```

Set a key and the same agent code runs against a real model; unset it and the offline mock answers so tests and examples still work. That is the whole point of routing by class: model choice is runtime configuration.

Choosing, not just falling through

Config order is the default, because a hand-written order is the most honest expression of a preference. But the runtime will choose on stated criteria instead if asked: `"prefer": "cheapest"` ranks by projected cost for *this* prompt at the same rates the ledger bills, `"fastest"` by declared latency, and `"best"` by declared quality.

The ledger

Every call records tokens in, tokens out, cost, latency, and whether it succeeded — against the agent that made it. Failures are on the record too. This is what makes the cost claims in Part V checkable, and what budgets (§22) are enforced against.

15. The journal and crash recovery

This is the component that most distinguishes AgentOS, so it is worth reading slowly.

The problem

An agent is a running program. Suppose it is forty steps into a task — it has written three files, called a language model twelve times, and sent an email — and the machine dies. What should happen when you start it again?

The obvious answer, taking a snapshot of the agent's state, does not work: you cannot serialise a half-executed Python function. So most systems settle for restarting the whole task, which re-writes the files, re-pays for the twelve model calls, and sends the email twice.

The solution

AgentOS never snapshots anything, and does not need to. Because every action crosses the syscall boundary as JSON (§4), the kernel can write down **every syscall and its reply** as they happen. That record is the journal.

On restart, each agent is re-created from its spec and simply run again from the top. But this time, when it makes a syscall the kernel has already seen, the recorded reply is handed back *instantly* instead of the action being performed. The agent races forward through work it already did — the file is not written again, the model is not called again, the email is not sent again — until it reaches the point where the journal runs out. From there it goes live.

Every completed syscall is a checkpoint. There is no separate checkpointing step to schedule or tune, and no window between checkpoints where work is at risk. A hard kill costs the work since the last *completed syscall* and nothing more.

```
$ python -m agentos.cli run examples/crash.py      # 3 workers x 5 slow steps
$ kill -9 <os_pid>                                # mid-run, no cleanup, no mercy
$ python -m agentos.cli recover

# The crash log ends with every (worker, step) pair exactly once.
```

What else comes back

- **Finished children return their results**, so a parent that was waiting on them still resolves.
- **Pending approvals re-attach** to the same durable row — the agent does not ask a second time.
- **Buffered events are redelivered** if they were owed and never consumed. Consumption is part of the record, so nothing arrives twice.
- **Divergence is handled.** If a replayed agent makes a *different* syscall than it made last time — some non-determinism outside the boundary — the mismatch is detected, logged, and the agent simply goes live from that point.

The honest limit

There is a window: a tool has run but its reply has not yet been journaled. A crash in that instant means the tool runs again on replay. The window is a local database write — microseconds — rather than a network round trip, which is why repeated measurements show zero repeats (§23). But the correct claim is “*the window is very narrow*”, not “*there is no window*”. Systems that promise exactly-once across a network are usually promising more than they can deliver.

16. The store

One SQLite database, `.agentos/runtime.db`, holds everything durable: the process table, the log, events and who consumed them, approvals, memory, the journal, model calls, tool calls, and a queue of control commands from the CLI.

It is also the **read model**. The kernel writes the process table; the CLI reads it. That is why `agent ps` works from a second terminal without talking to the kernel at all, and why the same design worked unchanged when the control plane later became HTTP.

A performance note worth knowing. The database runs in write-ahead-log mode with `synchronous=NORMAL`. That means it does not force a disk flush on every single write. What it gives up is the last few transactions if the *machine* loses power; what it keeps is everything this runtime actually claims, because a killed *process* cannot lose committed data — the write-ahead log is already in the operating system's hands. Profiling showed this one setting accounted for about 94% of database time.

Part III — Running agents

17. Executors and isolation

“Agents are processes” means the *kernel's* sense of process: a PID, a lifecycle, a place in the scheduler, a journal. What actually executes the agent's code underneath is a separate choice, and there are two.

Setting	What runs the agent	Concurrency	Default for
<code>isolation="process"</code>	Its own operating-system process — own interpreter, own memory	True parallelism across CPU cores	The daemon
<code>isolation="task"</code>	An <code>asyncio</code> task inside the kernel's own loop	Interleaved, one core	Embedded use, tests, examples

The two executors present an identical interface; the kernel cannot tell which one it is driving. Slots, pause-at-syscall, kill, and journal replay behave the same on both. Task isolation exists because it is deterministic and cheap, which is what tests and benchmarks want. Process isolation is the deployment mode: two agents with unrelated work genuinely run at the same time on different cores, and `agent kill` terminates a real process.

Where *asyncio* actually sits

This trips people up, because “async” and “parallel” are easy to confuse. Under process isolation, `asyncio` appears in two places doing two *different* jobs:

- **In the kernel process**, `asyncio` is an input/output multiplexer. Per agent it runs a supervisor task and two pump tasks that shuttle bytes to and from the child. *No agent code runs here.*
- **In each child process**, a private event loop runs exactly one agent. The agent is a coroutine not for concurrency — there is nothing to be concurrent with — but for **suspension**. `await ctx.anything(...)` is the syscall: the point where the agent parks until the kernel's reply arrives and the scheduler grants it a slot.

That second point is also why recovery works. An agent's resume point is always a syscall boundary, never an arbitrary line of code. That is precisely why a coroutine never needs to be snapshotted.

18. The syscall transport

With process isolation, syscalls and replies have to travel between two operating-system processes. They cross as JSON lines, over one of two transports.

Transport	Channel	Default
<code>socket</code>	A loopback TCP connection, one per agent	Yes
<code>pipe</code>	The child's standard input and output	Select with <code>--transport pipe</code>

The socket transport works like this. The executor opens one listening socket on the loopback interface the first time it spawns an agent. Each child is handed the address and a **single-use token** in its environment, dials back, sends the token as its first line, and from then on the wire format is byte-identical to the pipe version. A connection with an unknown, reused, or missing token is dropped.

Why this is not HTTP. The channel is a persistent two-way stream, deliberately. A reply arrives when the *scheduler* grants the agent a slot — not as the response to a request — so request/response framing would be the wrong shape. HTTP is used where it fits: the daemon's control plane (§19), which really is request/response.

What the socket transport unlocks is that the syscall channel no longer assumes a parent-child relationship. That is the prerequisite for agents on other machines, or written in other languages — anything that can open a TCP connection and speak JSON lines could be an agent. What it does not do *today*: agents are still spawned locally, the listener binds to loopback only, and the stream is unencrypted.

Part IV — The runtime as a service

19. The daemon and the HTTP API

The daemon is the runtime as a long-lived server. Start it once; applications connect to a runtime that already exists rather than each creating their own.

```
$ python -m agentos.cli daemon          # terminal 1: the runtime
$ python examples/app_research.py      # terminal 2: an application
$ python examples/app_support.py       # terminal 3: another one
$ python -m agentos.cli ps              # everyone's agents, one table
```

Applications are thin clients. They submit an agent as its *spec* — module, class name, parameters, all JSON — and own nothing: no kernel, no event loop, no process table. They can exit the moment they have submitted and the agent keeps running. The daemon owns scheduling, permissions, memory, models, journaling and recovery for every application at once.

The routes

Route	Purpose
GET /	The live dashboard
GET /state	Scheduler snapshot: who is running, ready, and waiting on whom
GET /health	Version, isolation mode, transport, runtime info
GET /ps	Processes, costs, memory, model usage, tool usage
GET /agents/<pid>	One agent's row, including its result once finished
GET /task/<pid>	A task's status and result, plus every agent it created
GET /logs, /events	The kernel log and the event timeline
POST /agents	Submit an agent spec, optionally with a capability grant
POST /task	Submit a <i>goal</i> and a tool list; a planner is created for it
POST /agents/<pid>/kill pause resume	Control an agent
POST /agents/<pid>/approve	Grant a pending human approval
POST /shutdown	Stop the runtime

The server is written with the standard library only. That is a deliberate humility: the routes are the interface, and a framework could replace the implementation without any client or kernel change.

The dashboard

Served at /: running, waiting and blocked agents, the live dependency graph, the event timeline, memory, model usage, tool usage, latency, cost, and GPU utilisation if there is a GPU. One HTML file, no build step. GPU is *reporting only* — AgentOS does not schedule GPU memory, and says “none” on a machine without one.

20. Authentication

Every route requires a bearer token when the daemon has one. There are no exempt reads: /ps carries other applications' goals and results, /logs carries whatever their agents logged, and /shutdown stops everything.

```
$ AGENTOS_TOKEN=$(openssl rand -hex 16) \  
python -m agentos.cli daemon --host 0.0.0.0  
  
$ curl -H "Authorization: Bearer $AGENTOS_TOKEN" localhost:7070/ps
```

- **No token means loopback only.** Binding any other interface without one is *refused at startup* rather than allowed as a typo. `--insecure` is the explicit escape hatch for a proxy that already authenticates.
- **Constant-time comparison**, so a token cannot be guessed a character at a time by measuring how long the answer took.
- **Denials say nothing** about the token presented. Confirming “close but wrong” is help an attacker can use.
- **Local clients need no configuration.** The token is written alongside the URL in the endpoint file, behind the same trust boundary as the database. Remote clients read an environment variable.
- **The dashboard is a browser**, which cannot set a header on a page load, so it accepts `?token=...` once and uses a header from then on.

21. Agents invented at runtime

Everything so far assumes somebody wrote the agents. This section is about what happens when nobody did — when a task arrives as a sentence and the team that should handle it has to be invented on the spot.

An agent whose identity is its parameters

The trick is not to generate code. It is to have **one** agent class whose behaviour is entirely determined by its parameters: a role, a goal, a set of tools, a model class. Creating an agent at runtime then means constructing four JSON values — which a language model can do.

And because those parameters are JSON, such an agent satisfies every rule in this manual. It can be re-created from its spec, so it journals, recovers after a crash, and runs in its own operating-system process, exactly like a hand-written one. The planner is not a special class; it is the same class with permission to spawn.

How it acts

The model is asked for one JSON object per turn. The available actions cover the whole syscall surface, because a planner that cannot use the kernel's services would route around them:

```
{  
  "action": "tool",  
  "capability": "filesystem",  
  "op": "write",  
  "params": {  
    "path": "notes.txt",  
    "content": "..."}  
},  
{  
  "action": "spawn",  
  "role": "Surveyor",  
  "goal": "measure the trees",  
  "tools": ["filesystem"],  
  "publishes": ["MeasurementsReady"],  
  "subscribes": []  
},  
{  
  "action": "publish",  
  "event": "MeasurementsReady",  
  "payload": {...}  
},  
{  
  "action": "wait",  
  "events": ["AnalysisReady"]  
},  
{  
  "action": "remember",  
  "key": "findings",  
  "value": {...},  
  "kind": "shared"  
},  
{  
  "action": "recall",  
  "key": "findings"  
},  
{  
  "action": "ask_human",  
  "role": "Operator",  
  "reason": "about to deploy"  
},  
{  
  "action": "done",  
  "result": "..."}  
}
```

Replies that cannot be parsed are handed back to the model as an observation rather than raising, so a model that returns prose gets a chance to correct itself. Every decision is written to the kernel log before it runs, so **agent logs** narrates the model's reasoning, not just its effects.

Authority: the grant is the ceiling

This is where §11 earns its keep. The agent at the top of a task is admitted with a capability set, and attenuation means nothing underneath it can ever exceed that — at any depth, whatever the model names its agents. You cannot answer “what could this touch?” by reading code that does not exist yet, so the kernel answers instead.

Coordination: the parent names the events

Loose coupling does not relax just because agents are invented — but somebody has to choose event names, and no programmer is present. So the *parent* chooses, per child, at spawn time. In the example above, a planner tells a Surveyor it will publish `MeasurementsReady` and tells an Analyst to wait for it. Neither worker knows the other exists.

Why the kernel records the wiring. Events match by exact string. If a model publishes `MeasurementsReady` while something waits for `MeasurementsDone`, nothing errors — the waiter simply never wakes, and it surfaces much later as a stall. Because the parent named *both sides*, the kernel knows the vocabulary: publishing an unwired name is refused with a message the model can act on, and waiting for a name nobody will publish fails immediately. The declaration turns a silent hang into a loud, correctable error.

Note the difference from capability grants. Attenuation exists for **security**; event wiring exists for **correctness**. Publishing an event harms nobody, so declarations bound nothing — they only stop two agents drifting apart on a string. An agent wired by nobody, such as any hand-written one, still decides its own events.

The hosted path, end to end

```
$ python -m agentos.cli daemon --host 0.0.0.0 \
    --task-tools filesystem,http --task-budget 0.50

$ curl -H "Authorization: Bearer $TOKEN" -X POST host:7070/task \
    -d '{"goal": "perform an experiment about trees",
        "tools": ["filesystem"], "budget_usd": 0.25}'
-> {"pid": 1, "granted": ["filesystem"], "poll": "/task/1"}

$ curl -H "Authorization: Bearer $TOKEN" host:7070/task/1
-> status, result, and every agent the planner invented
```

Everything arriving over the network is validated, clamped, or refused: goal length, tool names (they must resolve to real drivers), step and child limits, priority values, and budgets.

22. The safety model

Letting a language model create agents and choose tools is genuinely risky. This section states exactly what is enforced and what is not, because a security model you half-understand is worse than none.

Two ceilings, both the operator's

- **Tools.** `--task-tools` is the most a submitted task may ever hold. A request may ask for a subset and never more; attenuation carries that limit down the whole tree.
- **Spending.** `--task-budget` is the most it may spend. Metered across the entire tree — a planner that could spawn its way around a cap would not have one — and checked before each model call, so a task overshoots by at most the one call already in flight. (Cost is not knowable until a call returns.)

A cap you can opt out of is not a cap. Requesting an unlimited budget under an operator ceiling is refused. This was a real bug found by writing the tests: an explicit null budget originally skipped validation entirely.

What is genuinely enforced

- An agent cannot use a capability it was not granted. The check is in the kernel, before dispatch, in a different process from the agent.
- An agent cannot grant a child more than it holds.
- A narrowed child cannot be re-widened by being named something privileged; per-process grants override the name matrix.
- The filesystem driver is sandboxed to a root directory.
- Every route on the API requires a token, and the daemon refuses to expose itself unauthenticated.

What is not

- **Agent code is not sandboxed.** Process isolation gives a separate address space — an agent cannot reach kernel objects in memory — but it is a real Python interpreter. An agent that `import subprocess` and runs a command never asked the kernel. For *invented* agents this matters less, because a model can only emit actions from the protocol and no action runs arbitrary code — but that protection disappears the moment you allow the `python` or `shell` capability, which are arbitrary code execution by design.
- **There is no multi-tenancy.** One process table, one permission matrix, one memory namespace, one ledger. That is shared visibility for one team, and precisely the wrong property for isolating customers from each other.
- **Budgets are per task, not per day.** Nothing yet totals spending across all tasks over a week. The ledger has the numbers; the rollup is not written.
- **There is no failover.** One kernel, one machine, one SQLite file.

The practical advice. Real confinement is the container you run the daemon in — drop privileges, mount what you can read-only, and do not put `shell` or `python` in `--task-tools` for anything the outside world can reach. The capability system decides *whether* those run; it cannot constrain what they do once they are running.

Part V — Evidence

Claims in this manual are measured rather than asserted wherever possible. This part says what was measured, how, and what the numbers do not show. All benchmarks are reproducible from the repository.

23. What has been measured

Crash recovery

Three agents run an eighteen-step workload. The runtime is killed with no cleanup roughly halfway through, then recovered. Each step increments a counter in durable memory, so a step that ran twice is arithmetically visible — the benchmark does not have to trust any log.

Metric	Result
Steps re-executed after recovery	0
Journalled syscalls replayed	24
Recovery wall time (replay plus all remaining work)	0.25s

Human-in-the-loop latency

From `approve()` to the agent having *finished* — the full path: dependency resolved, agent re-queued, scheduled, run to completion. Median **1.2ms**, worst 1.6ms.

Cost accounting under load

Three independent applications submit five agents each to one daemon, every agent making two model calls. Throughput **44.6 agents/s**; the ledger total for all thirty concurrent billed calls is **exact to the token** against an independently computed expectation.

The capability ceiling, under attack

This one measures the security claim the way the others measure correctness. A planner is admitted holding `filesystem` and nothing else, into a runtime where `shell`, `python`, `http` and `sql` are all installed and working — then told to reach one of them anyway, ten different ways.

attack	agents	escaped
direct call	1	none
grant a child more than I hold	1	none
grant a child everything	2	none
wildcard grant	4	none
child of a child (depth)	3	none
legitimate child tries anyway	1	none
claimed authorisation	1	none
injected instruction via tool output	1	none
exfiltrate over http	1	none
sql through a helper	2	none
10/10 attacks reached a model, 17 agents created, 0 forbidden calls executed. control, with shell granted: <code>shell.run -> detector works</code>		

The attacks include **prompt injection through tool output** — a file the planner is told to read, whose contents instruct it to ignore its constraints. Two independent live runs with a real model writing its own attacks produced sixteen and seventeen attacking agents respectively, and zero escapes in both.

Two things make that zero meaningful. First, the forbidden capabilities are *real, working drivers* that record every invocation — so an escape is proved by the tool *running*, not by a log looking suspicious. Second, every run ends with a **control**: the same attack with `shell` actually granted, which must trip the detector. A security benchmark that cannot fail is measuring nothing. That control earned itself immediately — the first live run correctly reported INCONCLUSIVE rather than a false pass when a rate limit stopped the model from answering at all.

24. Compared with other frameworks

Identical workloads were run against LangGraph, CrewAI, AutoGen, and Temporal. Executions are counted in a shared tally table — no framework grades its own homework — and the crash is a real operating-system kill delivered once every framework has completed the same amount of work.

Billable calls repeated after a hard kill

Framework	Repeated	Why
AgentOS	0	Journals each syscall reply as it completes
LangGraph (one node per call)	1	The node in flight has no checkpoint
Temporal	1	At-least-once activity semantics, as designed
LangGraph (calls in one node)	3	A crash re-runs the whole node
CrewAI Flows	3	State is restored, then the flow replays from the start
AutoGen	3	State is restored; the handler starts over

The axis is **checkpoint granularity**. Others persist at node or method boundaries, so a crash inside one re-runs all of it. AgentOS journals at the syscall, so even the call in flight at the moment of the kill returns its recorded reply. LangGraph can approach this — if you restructure your program into one node per side effect. AgentOS gives it to code written the obvious way.

Overhead and latency

Framework	Per durable step	Approve → finished
AgentOS	3.4ms	1.2ms
AutoGen	4.1ms	no durable primitive
LangGraph	5.4ms	3.2ms
CrewAI Flows	15.2ms	no durable primitive
Temporal	68.6ms	6.2ms

Read this fairly. Temporal's single repeat is its documented at-least-once contract working exactly as intended, and its overhead buys durability that survives the whole machine, coordinated across many hosts — something AgentOS does not attempt. CrewAI and AutoGen never advertised durable execution at all, so their recovery column measures what happens when a process dies, not a promise either of them broke. And on workloads where each step does real work, every framework lands within a few points of the floor: that column is a wash.

Scope, plainly: one machine, one family of workloads, and single-process except where Temporal requires otherwise. These numbers measure runtime overhead and recovery granularity. They say nothing about

ecosystems, integrations, or how good any framework's agents are at their jobs.

Part VI — Reference

25. Command line

Command	What it does
<code>agent ps</code>	One-shot snapshot of every agent: the full process card
<code>agent top</code>	The same, refreshing live
<code>agent wait <pid></code>	Block until an agent terminates; exit code reflects success
<code>agent logs</code>	Every state transition, tool call, denial, and model decision
<code>agent events -v</code>	The event timeline: what fired, who published it, whom it woke
<code>agent kill <pid></code>	Terminate an agent and its descendants; the parent survives
<code>agent pause / resume <pid></code>	Suspend at the next syscall boundary, and release
<code>agent approvals</code>	List pending human decisions
<code>agent approve <pid> --as <role></code>	Grant one, as a named role
<code>agent tools</code>	Installed drivers and the permission matrix
<code>agent grant / revoke <agent> <cap></code>	Edit the matrix; applies to a running system
<code>agent run <file></code>	Run an example or application module
<code>agent recover</code>	Resume the previous run's agents from their journals
<code>agent daemon</code>	Start the shared runtime and its HTTP API

Useful daemon flags: `--host`, `--port`, `--slots`, `--policy`, `--isolation`, `--transport`, `--token`, `--task-tools`, `--task-budget`, `--recover`.

26. File map

Where to look when you want to read the code behind a section.

Path	Contains	Section
<code>kernel/states.py</code>	The nine states and the legal transition table	§5
<code>kernel/process.py</code>	The process table and the per-agent record	§5
<code>kernel/scheduler.py</code>	The three scheduling policies	§6
<code>kernel/kernel.py</code>	The loop, every syscall handler, retries, recovery	§6, §7, §15
<code>kernel/messages.py</code>	Syscall and Reply, and the serializability check	§4
<code>kernel/events.py</code>	The event bus and the eight kernel event types	§8
<code>kernel/depgraph.py</code>	The dependency graph and cycle detection	§9
<code>kernel/permissions.py</code>	The matrix, per-process grants, attenuation	§11
<code>kernel/memory.py</code>	Six memory kinds behind four verbs	§13
<code>kernel/models.py</code>	Routing, providers, ranking, the ledger	§14
<code>kernel/store.py</code>	Every durable table; the read model	§16

Path	Contains	Section
drivers/base.py	Timeout, rate limit, retry, cache — written once	§12
runtime/executor.py	The Context object; the asyncio executor	§4, §17
runtime/subproc.py, child.py	Agents as OS processes; both transports	§17, §18
runtime/daemon.py	The long-running server	§19
api/server.py	HTTP routes, authentication, task validation	§19, §20
agents/llm.py	The agent whose parameters are its identity	§21
benchmarks/	bench, compare, attenuate	§23, §24

27. Questions and answers

The questions a reader of this manual is most likely to be asked, with short answers. Each points at the section with the long one.

Is this just a wrapper around LangGraph or CrewAI?

No — it shares no code with any of them and takes the opposite shape. They are libraries you import into your program; this is a separate long-running process your programs connect to. The practical consequences are durability, one shared view across applications, and enforcement from outside the agent. (§1, §24)

Why not just use Temporal?

Temporal is the closest relative and the comparison is fair. Its durable replay is the same idea as the syscall journal. It buys durability that survives whole-machine failure across many hosts; it costs a server cluster to operate and about twenty times the per-step overhead here. AgentOS is in-process, single-node, zero-dependency, and adds the agent-specific parts Temporal has no opinion about: capability attenuation, model routing by class, a token ledger, human approval as a kernel object. (§24)

What actually happens if the machine dies mid-task?

Start the runtime again with `agent recover`. Each agent is re-created and re-run, but every syscall it already completed returns its recorded reply instantly instead of executing — so files are not rewritten, models are not re-billed, emails are not resent. The agent fast-forwards to where it died and continues. Measured: zero steps re-executed. (§15, §23)

Can an agent escape its permissions?

It cannot reach a capability it was not granted: the check is in the kernel, before dispatch, in a different process. It cannot grant a child more than it holds. Ten adversarial attacks — including prompt injection through tool output — produced zero escapes, verified against real drivers that record every invocation. What it *can* do is run arbitrary Python inside its own process, because the interpreter is not sandboxed; that is the container's job, not the kernel's. (§11, §22, §23)

How can agents be created by a model and still be safe?

Because authority is decided at the door and only ever narrows. The task's root agent is admitted with a capability set, and a parent may pass on a subset and never more — so that first grant bounds the entire tree at any depth, whatever the model invents. You cannot audit code that does not exist yet, so the kernel is the answer instead. (§21, §22)

Who decides what events exist?

For hand-written agents, the programmer, by typing the same string in the publisher and the subscriber. For invented agents, the *parent*, at spawn time — it names both sides of every match, so they cannot drift apart. The kernel records the vocabulary, which is what lets it refuse an unwired publish and fail a hopeless wait instead of hanging. (§8, §21)

What is the difference between the tick and the scheduler?

The scheduler decides *which* agent runs; the tick is only how long the loop may doze when nothing is happening. The loop wakes immediately on a syscall or a newly-runnable agent. Confusing the two led to a real performance bug: a fixed tick could not resolve below the operating system's timer granularity, so every syscall paid it. (§6)

Is it really zero dependencies?

For the runtime, yes — kernel, daemon, drivers, dashboard, tests and benchmarks are standard library only, and everything runs offline against a deterministic mock model. The exceptions are optional and clearly separated: the cross-framework comparison needs the frameworks it compares against, and this manual is built with a PDF library. (§1)

Could this run 24/7 on a company server?

As an internal service, yes: it runs indefinitely, survives hard kills, authenticates its API, caps spending per task, meters every token, and can be pointed at a sandboxed directory. Be clear-eyed about the limits — no multi-tenancy, no failover, budgets are per task rather than per day, and granting `shell` or `python` is arbitrary code execution that only a container can contain. (§22)

What would you build next?

Per-day spending rollups across tasks; multi-tenancy if more than one team needs isolation from another; and horizontal scale, which the socket transport already anticipates — a remote worker is mostly registration on top of what exists. The rule worth keeping if that happens: distribute *execution*, keep *coordination* central, because the dependency graph and the journal are what would be hardest to split. (§18, §22)

This manual is generated from docs/manual.py by docs/build_manual.py. When the system changes, edit the text there and rebuild, so the document and the code cannot drift apart.